

Fantasy Hockey Auction Draft Strategy

A Project on Solving Fantasy Hockey Auction Drafts

Problem Statement and Motivation

- **Objective:** Identify an optimal strategy for fantasy hockey auction drafts.
- **Game Structure:**
 - **Players:** n players draft a team of:
 - o forwards
 - d defensemen
 - g goalies
 - **Budget:** Each player has a fixed budget of b .
 - **Auction Format:**
 - English auction—each round:
 - Players **raise** (bid increases by 1) or **fold** (exit the round).
 - The last remaining player wins the athlete at their bid.
 - **Constraints:** Budget management and team composition rules.
- **Assumptions:** All athlete values are known beforehand.
- **Motivation:** Inspired by my personal yearly fantasy hockey auction league—an engaging and strategic game.

Approach 1: AlphaZero Method

Overview

- Leverages self-play and Monte Carlo Tree Search (MCTS).
- Operates in a **perfect information** setting.
- **Auction Dynamics:**
 - In each round, players choose to **raise** or **fold**.
 - A raise increases the bid by 1; folding exits the round.
 - The last remaining player wins the athlete at their bid.

Approach 1: AlphaZero Method

Mathematical Formulation

For a given game state (s):

$$a^* = \arg \max_{a \in \{\text{raise}, \text{fold}\}} \pi(s, a)$$

- $\pi(s, a)$: Policy function learned during self-play.

Approach 1: AlphaZero Method

Challenges

- **Enormous Search Space:**
 - With n players and budget b , the number of possible turns is huge.
 - The game tree complexity grows as approximately $O(2^{nb})$ (compared to 60–80 moves in chess).
- **Computational Overhead:**

MCTS becomes expensive as the game tree expands.

Approach 2: PPO (Proximal Policy Optimization)

Overview & How It Works

- An actor-critic policy gradient method.
- Uses a **clipped surrogate objective** to ensure stable policy updates.
- **Auction Modeling:**
 - Each round is approximated as a Vickrey auction.
 - Bids are modeled as a percentage of the remaining budget.

Approach 2: PPO

Mathematical Optimization Decision

- **Surrogate Objective:**

$$L^{CLIP}(\theta) = \mathbb{E}_t \left[\min (r_t(\theta) A_t, \text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon) A_t) \right]$$

where:

- $r_t(\theta) = \frac{\pi_\theta(a_t|s_t)}{\pi_{\theta_{\text{old}}}(a_t|s_t)}$
- A_t is the advantage estimate.
- ϵ is a small hyperparameter.

- **Bid Modeling:**

The bid is given by:

$$\text{bid} = b_r \cdot x, \quad x \sim \text{Beta}(\alpha, \beta)$$

where b_r is the maximum bid we can make, and α, β are generated from mean, variance outputs of our policy network.

Approach 2: PPO

Challenges

- **Credit Assignment:**
Early moves have a larger impact on the final outcome, but PPO's structure treats all moves similarly.
- **Imperfect Information:**
The auction's uncertainty complicates training. It is also unclear that a pure strategy is optimal.

Approach 3: Deep CFR (Counterfactual Regret Minimization)

Overview & How It Works

- Uses regret minimization to approach Nash equilibria in imperfect information games.
- **Key Distinction:**
Instead of directly predicting regret, our network predicts the **expected reward** $Q(s, a)$ for each action a at state s .
- **Auction Modeling:**
 - Discretize the action space as percentages of the remaining budget (0% to 100%).

Approach 3: Deep CFR

Revised Mathematical Formulation

- **Expected Reward Prediction:**

For each state s and action a :

$$Q(s, a) = \text{Expected Reward given action } a$$

- **Regret Calculation:**

Compute regret relative to the expected value of the current policy $\sigma(s)$:

$$R(s, a) = Q(s, a) - \sum_{a'} \sigma(s, a') Q(s, a')$$

- **Policy Update via Regret Matching:**

Play the policy based on positive regrets:

$$\sigma(s, a) = \frac{\max(R(s, a), 0)}{\sum_{a'} \max(R(s, a'), 0)}$$

Approach 3: Deep CFR

Neural Network Approximation & Special Notes

- **Network Optimization:**

Train the network to approximate $Q(s, a)$ by minimizing:

$$\min_{\theta} \mathbb{E}_{s \sim \mathcal{D}} \left[\sum_a \left(Q(s, a) - \hat{Q}(s, a; \theta) \right)^2 \right]$$

where $\hat{Q}(s, a; \theta)$ is the network's estimate.

- **Implementation Notes:**

- Start with a 2-player scenario to leverage stronger theoretical guarantees.
- Adjustments to the original Deep CFR formulation improve practical performance.
- The athlete pool is fixed per season—requiring retraining with updated projections.
- We train on all samples we have seen across all iterations.

Approach 3: Deep CFR

Policy Network Training with JSD Minimization

- **Additional Training:**
 - Train a separate policy network to capture the aggregate strategy.
- **Jensen-Shannon Divergence Minimization:**
 - **Objective:** Align the network's predicted distribution $\pi_{\phi}(s)$ with the aggregate policy $\pi_{\text{agg}}(s)$ observed across all iterations.
 - **Formulation:**

$$\min_{\phi} D_{JS} \left(\pi_{\text{agg}}(s) \parallel \pi_{\phi}(s) \right)$$

- **Benefit:**
 - Encourages the policy network to learn from the best strategies played so far.
 - Provides a robust policy approximation that generalizes across iterations.
 - Jensen-Shannon over KL-Divergence gave more stable outputs.

Approach 3: Deep CFR: Simplified Game Iteration 1

Final Means: [300, 295] , Final Budgets: [8, 3]

Value	Price	Winner	Bids	Curr. Budgets
115	60.0	1	[60, 69]	[100, 40]
110	37.0	0	[69, 37]	[63, 40]
80	37.0	0	[59, 37]	[26, 40]
75	24.0	1	[24, 37]	[26, 16]
50	12.0	0	[24, 12]	[14, 16]
45	0.0	1	[0, 13]	[14, 16]
40	13.0	1	[13, 15]	[14, 3]
35	3.0	0	[13, 3]	[11, 3]
25	3.0	0	[11, 3]	[8, 3]
20	0.0	1	[0, 3]	[8, 3]

Approach 3: Deep CFR: Simplified Game Iteration 3

Final Means: [325, 270] , Final Budgets: [29, 8]

Value	Price	Winner	Bids	Curr. Budgets
115	47.0	0	[88, 47]	[53, 100]
110	50.0	1	[50, 60]	[53, 50]
80	42.0	1	[42, 47]	[53, 8]
75	6.0	0	[50, 6]	[47, 8]
50	6.0	0	[35, 6]	[41, 8]
45	6.0	0	[40, 6]	[35, 8]
40	6.0	0	[35, 6]	[29, 8]
35	0.0	1	[0, 6]	[29, 8]
25	0.0	1	[0, 3]	[29, 8]
20	0.0	1	[0, 8]	[29, 8]

Approach 3: Deep CFR: Simplified Game Iteration 25

Final Means: [310, 285] , Final Budgets: [5, 1]

Value	Price	Winner	Bids	Curr. Budgets
115	61.0	1	[61, 70]	[100, 39]
110	36.0	0	[87, 36]	[64, 39]
80	36.0	0	[61, 36]	[28, 39]
75	26.0	1	[26, 36]	[28, 13]
50	11.0	0	[26, 11]	[17, 13]
45	11.0	0	[16, 11]	[6, 13]
40	6.0	1	[6, 11]	[6, 7]
35	6.0	1	[6, 6]	[6, 1]
25	1.0	0	[6, 1]	[5, 1]
20	0.0	1	[0, 1]	[5, 1]

Approach 3: Deep CFR: Simplified Game Iteration 175

Final Means: [300, 295] , Final Budgets: [0, 5]

Value	Price	Winner	Bids	Curr. Budgets
115	51.0	0	[58, 51]	[49, 100]
110	46.0	1	[46, 70]	[49, 54]
80	31.0	1	[31, 33]	[49, 23]
75	21.0	0	[46, 21]	[28, 23]
50	18.0	0	[26, 18]	[10, 23]
45	9.0	1	[9, 21]	[10, 14]
40	9.0	1	[9, 13]	[10, 5]
35	5.0	0	[9, 5]	[5, 5]
25	5.0	0	[5, 5]	[0, 5]
20	0.0	1	[0, 5]	[0, 5]

Conclusion

- **Summary:**
 - Explored advanced methods: AlphaZero, PPO, and Deep CFR.
 - Each approach addresses unique aspects of the fantasy hockey auction challenge.
- **Looking Ahead:**
 - Continuous refinement and experimental validation will bring us closer to an optimal bidding strategy.
 - Open to feedback and further discussion on these methods.